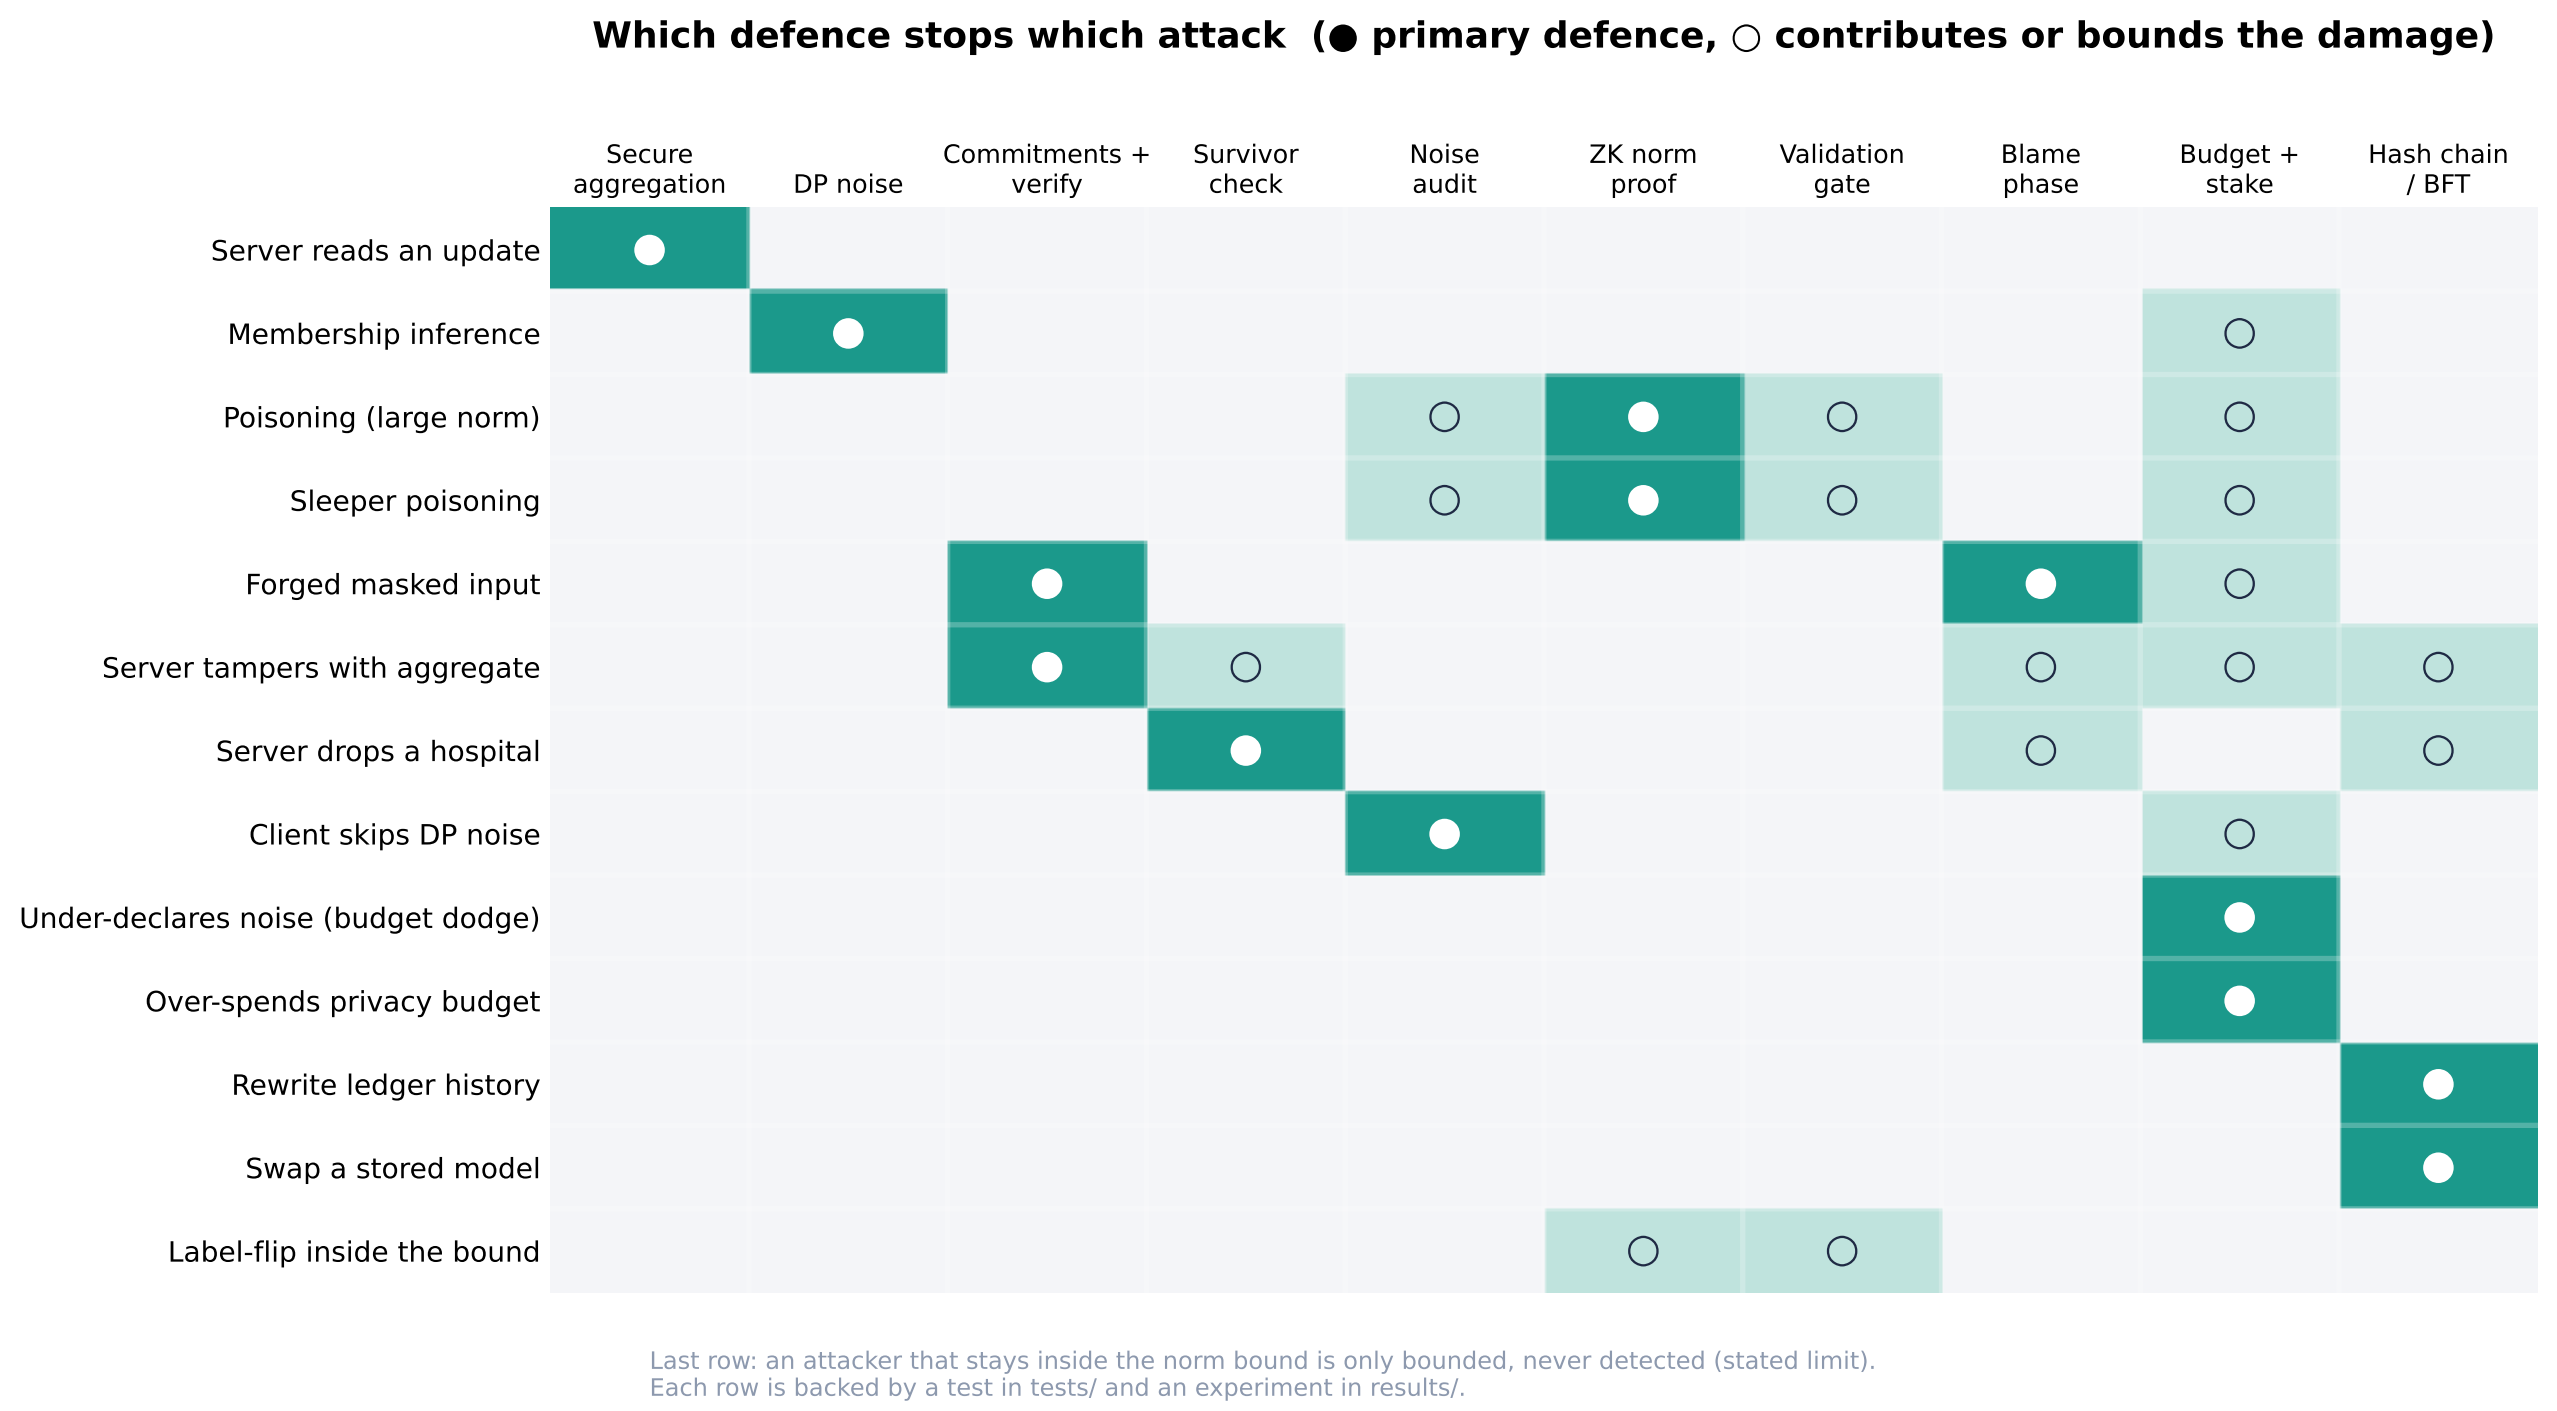
Which defence stops which attack (● primary defence, ○ contributes or bounds the damage)
Secure aggregation
DP noise
Commitments + verify
Survivor check
Noise audit
ZK norm proof
Validation gate
Blame phase
Budget + stake
Hash chain / BFT
Server reads an update
Membership inference
Poisoning (large norm)
Sleeper poisoning
Forged masked input
Server tampers with aggregate
Server drops a hospital
Client skips DP noise
Under-declares noise (budget dodge)
Over-spends privacy budget
Rewrite ledger history
Swap a stored model
Label-flip inside the bound
Last row: an attacker that stays inside the norm bound is only bounded, never detected (stated limit).
Each row is backed by a test in tests/ and an experiment in results/.